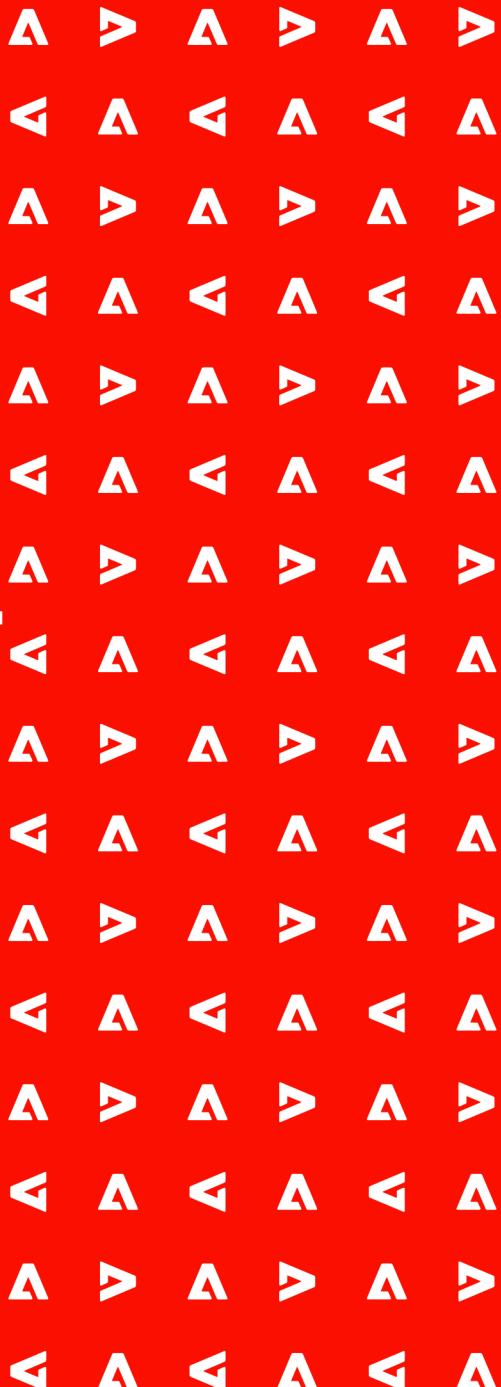




Fully Convolutional Mesh Autoencoder using Spatially Varying Kernels

Presenter: Yi Zhou

09/2020

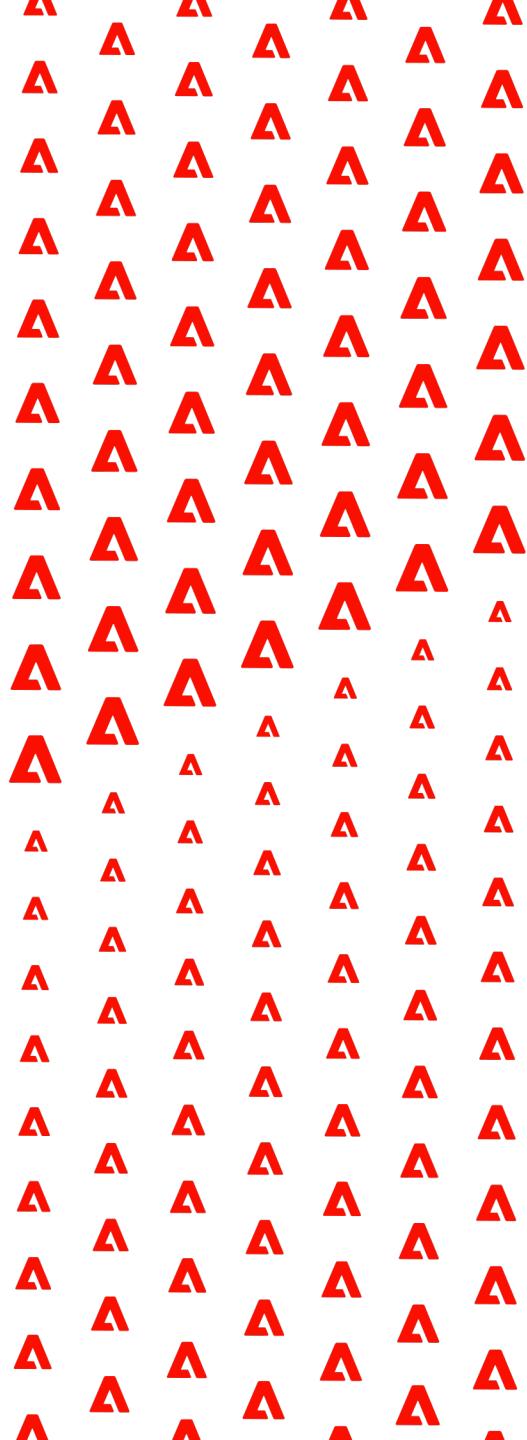


A Collaboration with [Facebook Reality Labs](#) and [USC](#):

Zhou Yi, [Chenglei Wu](#), [Zimo Li](#), [Chen Cao](#), [Yuting Ye](#), [Jason Saragih](#), [Hao Li](#), and [Yaser Sheikh](#).

"Fully Convolutional Mesh Autoencoder using Efficient Spatially Varying Kernels."

Accepted by Neurips 2020



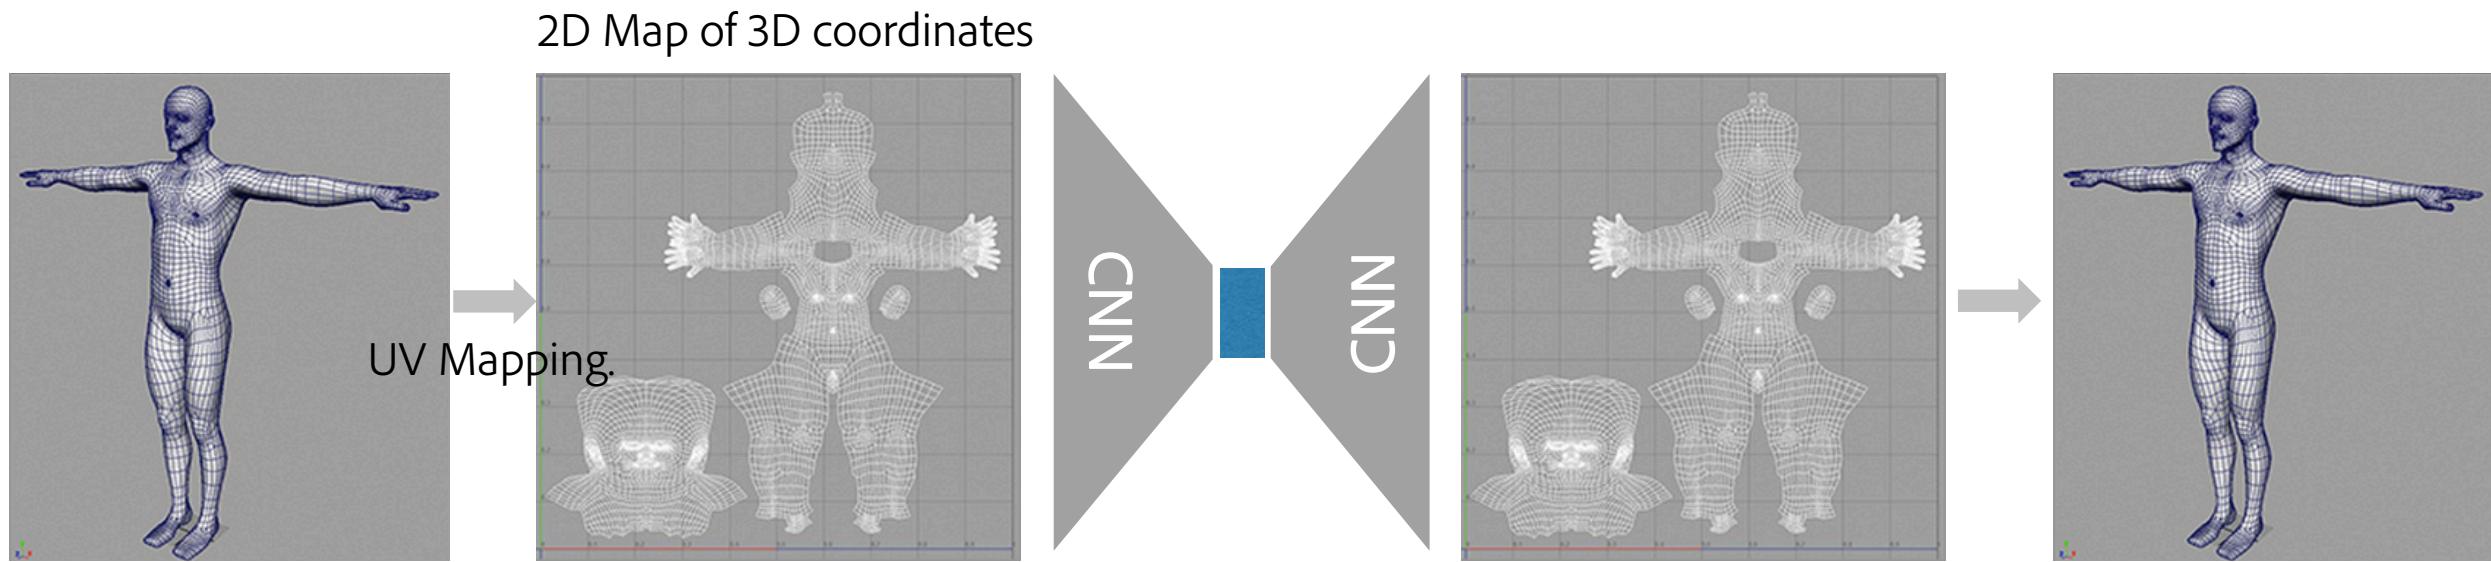
How to apply CNN on registered 3D meshes?

Registered Mesh: Mesh with the same number and order of vertices and edges.



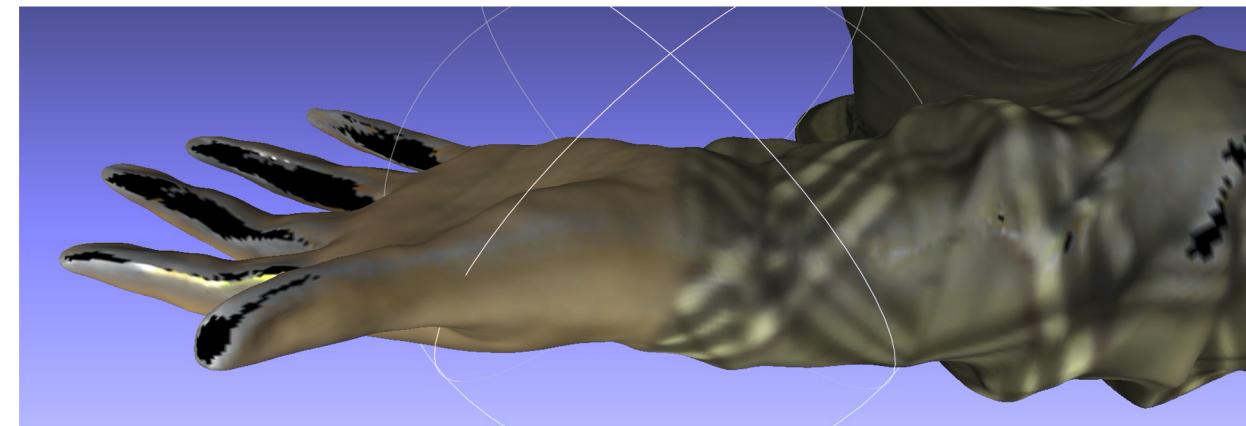
<http://dfaust.is.tue.mpg.de/>

Common Practice: 2D CNN in UV space

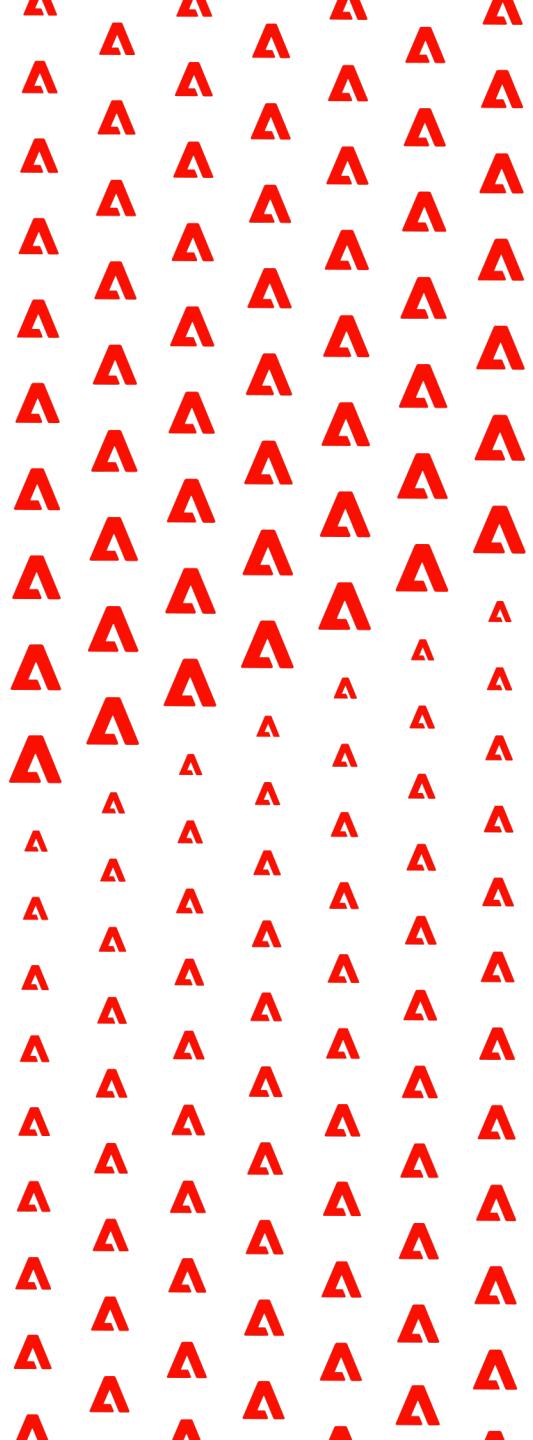


Problems:

1. Artifacts along seam lines and from distortion
2. Poor performance when reconstructing global deformation



CNN on 3D Meshes

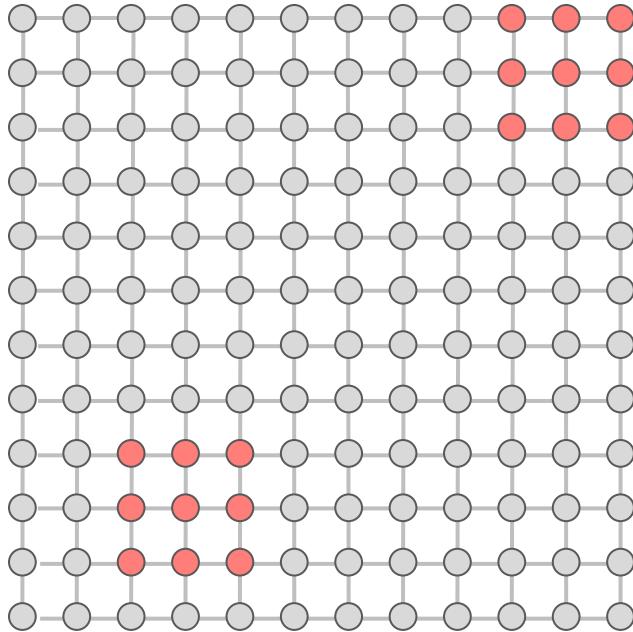


Difficulties

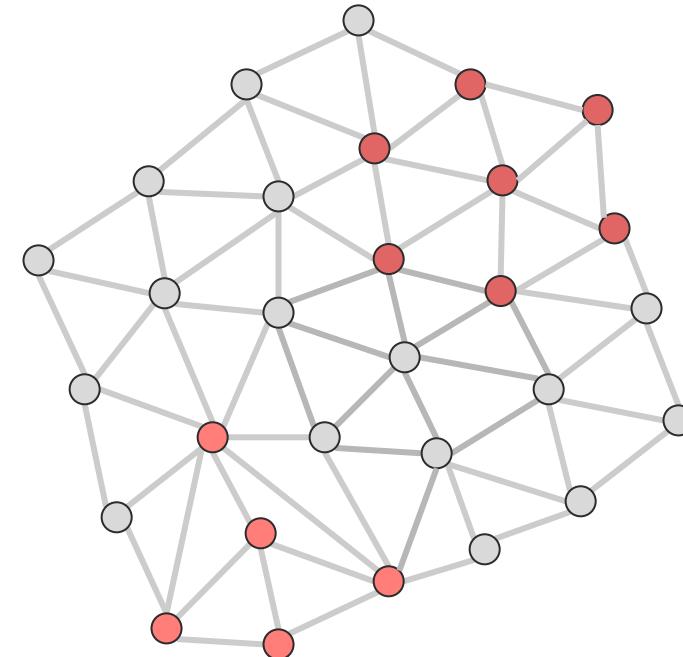
A Mesh is usually a non-uniform discretization.



Cannot sample uniform kernels on a non-uniform mesh



Shift-invariance Grid



Shift-variance Mesh

Existing Methods

Spectral Method:

(Chebyshev ...)

Lose Fidelity, unstable

Special Method:

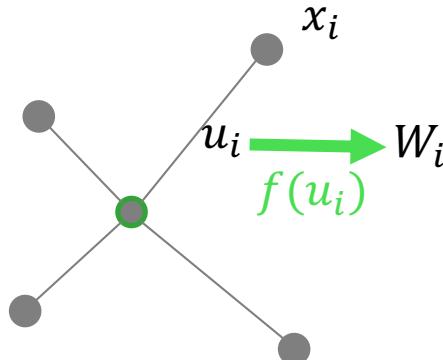
(Spiral CNN [Neural3DMM 2020])

Ad-hoc, limited to 2-D manifold

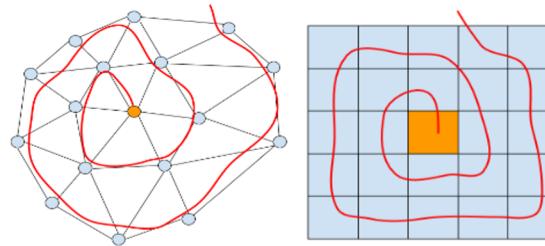
Feature-Conditioned Method:

(GAT, MoNet, FeaStNet ...)

Sensitive to big variations.

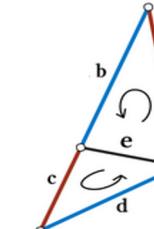


$$y = \sum W_i Ax_i$$



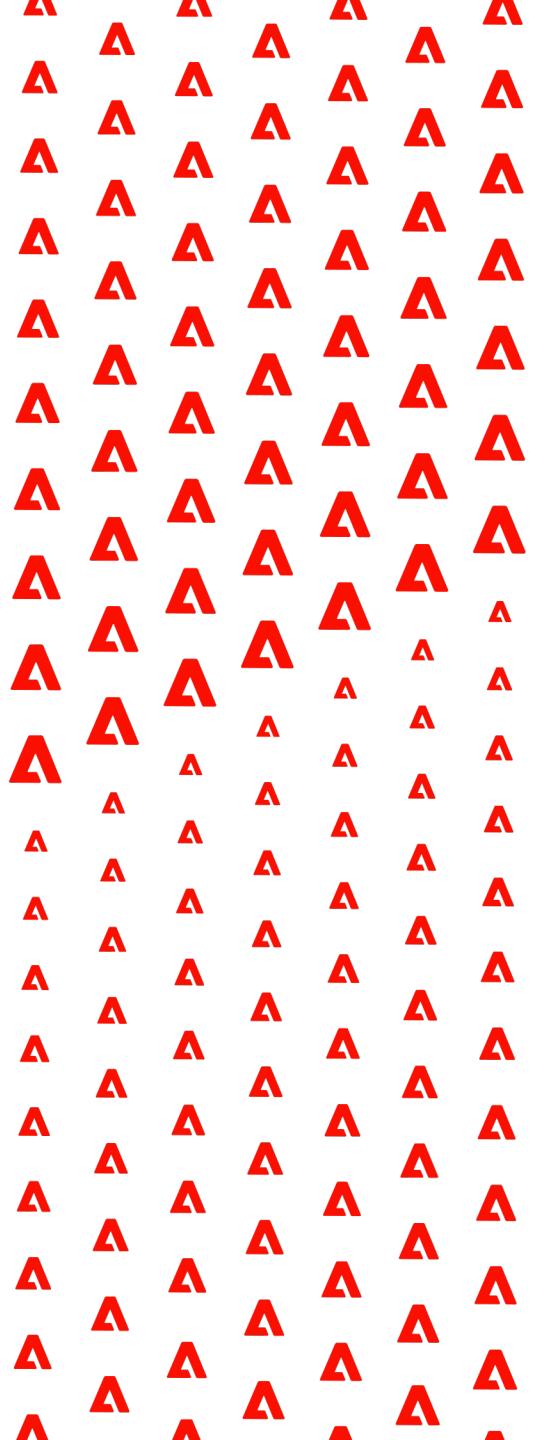
(EdgeCNN [MeshCNN 2019])

Slow, limited to 2-D manifold, Lose geometry information.



Mesh Convolution

Mesh CNN with Up and Down-Scaling



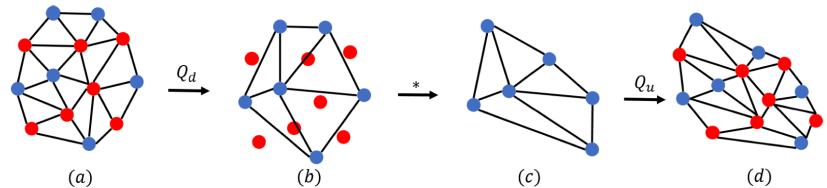
Existing Methods

Quadric Mesh Simplification:

[CoMA 2018, Neural3DMM 2019]

Fixed parameters

Overfit to one template mesh



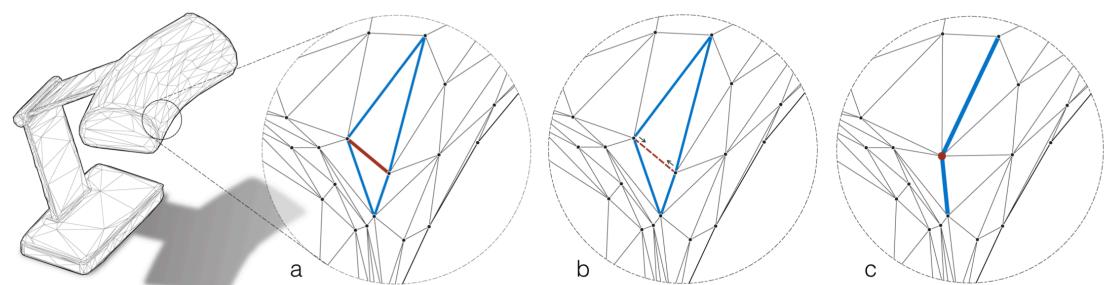
Dynamic Edge Collapsing:

[MeshCNN 2019]

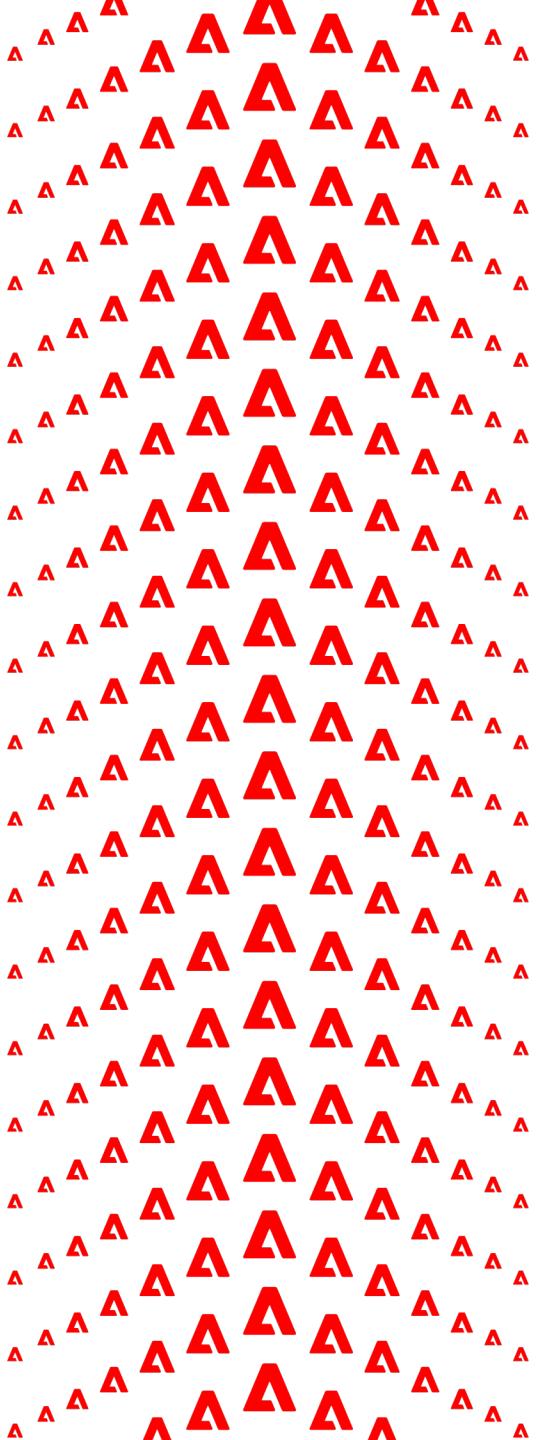
Very slow

limited to 2-D manifold

minimal downscaling size requirements.



Our Method

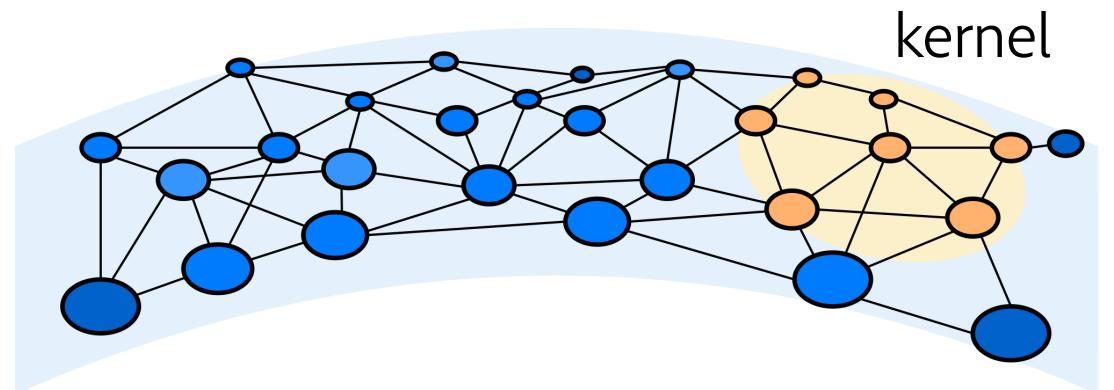


Insights

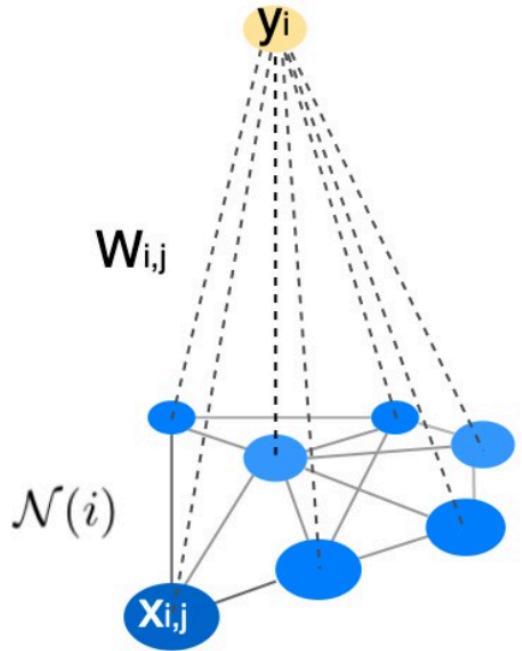
A continuous kernel can be shared in a continuous space.

A discretized kernel can be sampled from a continuous kernel.

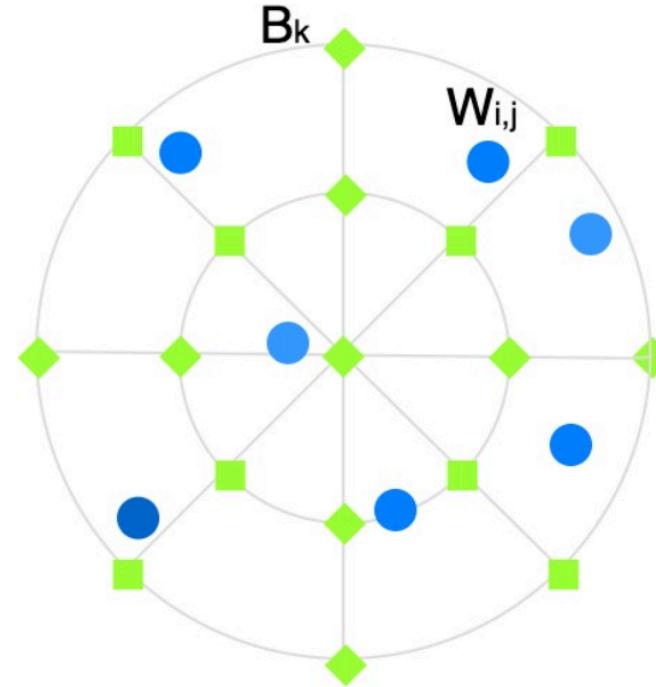
The sampling function needs to be defined per vertex locally.



Insights



Discrete Conv Filter at A Local Patch



- ◆ : Weight Basis on an imaginary grid
- : final weights

Our Convolution Operation

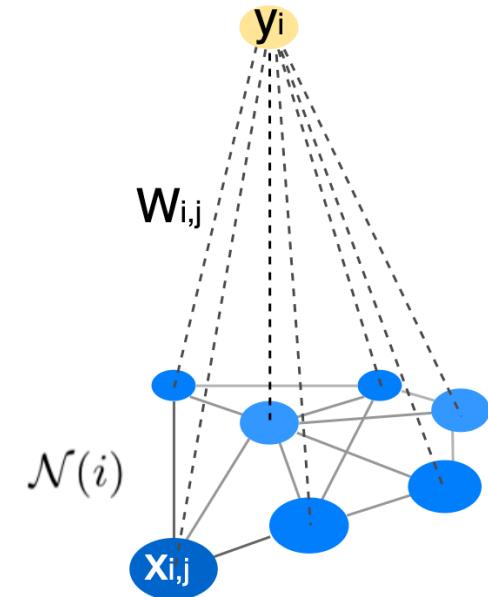
Each Convolution Layer has one kernel basis $B = \{\mathbf{B}_k\}_{k=1}^M, \mathbf{B}_k \in \mathbb{R}^{I \times O}$

Each edge j for a local vertex i has coefficients $A_{i,j} = \{\alpha_{i,j,k}\}_{k=1}^M, \alpha \in \mathbb{R}$:

The weight $\mathbf{W}_{i,j}$ on each edge is computed as $\mathbf{W}_{i,j} = \sum_{k=1}^M \alpha_{i,j,k} \mathbf{B}_k$

The output feature is computed as $\mathbf{y}_i = \sum_{x_{i,j} \in \mathcal{N}(i)} \mathbf{W}_{i,j} \mathbf{x}_{i,j} + \mathbf{b}$

B and $A_{i,j}$ are training parameters, shared across the dataset.



Our Pooling Operation

Observation: local density is non-uniform

Solution: Monte Carlo Integration with learned density coefficients

Formulation:

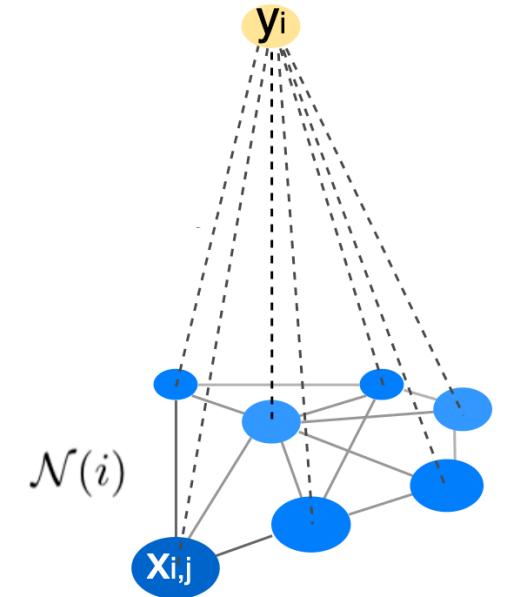
Each local vertex j has a density coefficient

$$\rho'_{i,j} = \frac{|\rho_{i,j}|}{\sum_{j=1}^{E_i} |\rho_{i,j}|}$$

The output feature is computed as

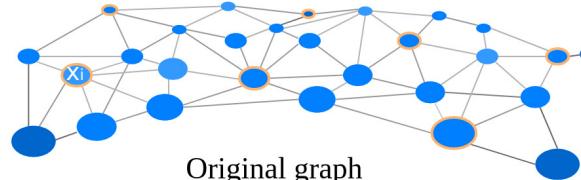
$$\mathbf{y}_i = \sum_{j \in \mathcal{N}(i)} \rho'_{i,j} \mathbf{x}_{i,j}$$

$|\rho_{i,j}|$ are training parameters, shared across the dataset.



Down and Up Scaling Based Only on Graph Connectivity

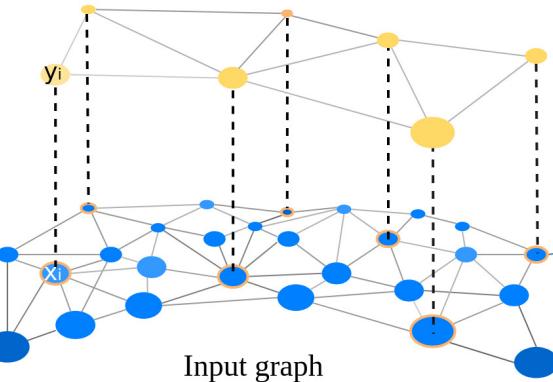
Down-sampling



Original graph

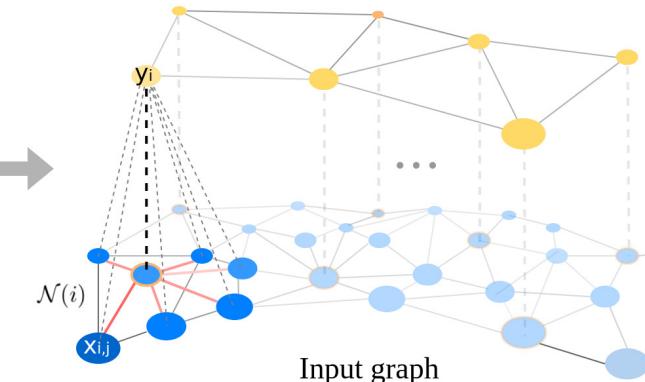
Select sampled vertices with stride=1.

Output graph



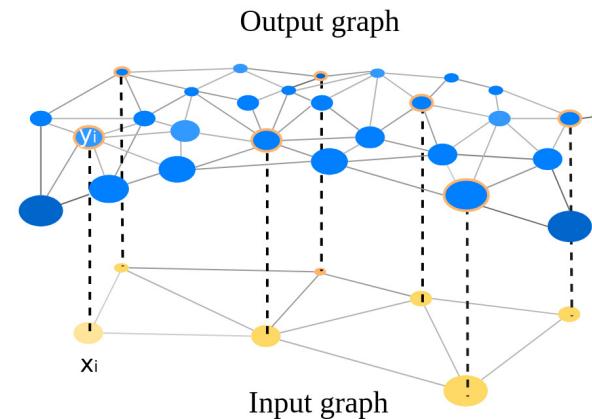
Create input/output graphs for Conv/Pool.

Output graph



Create edges between input and output graphs.

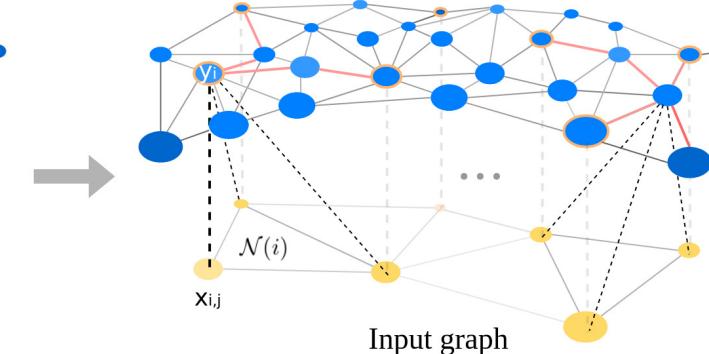
Up-sampling



Reverse

Output graph

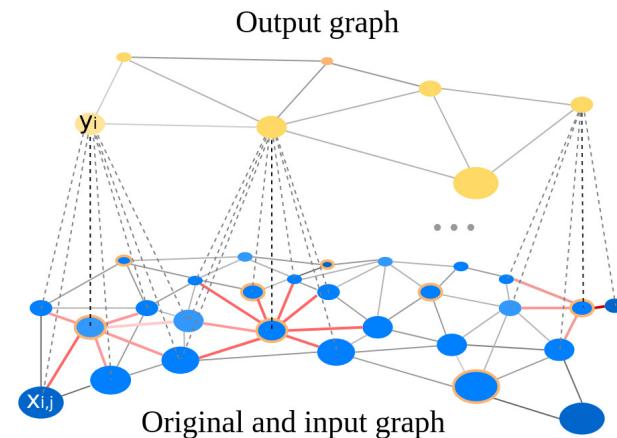
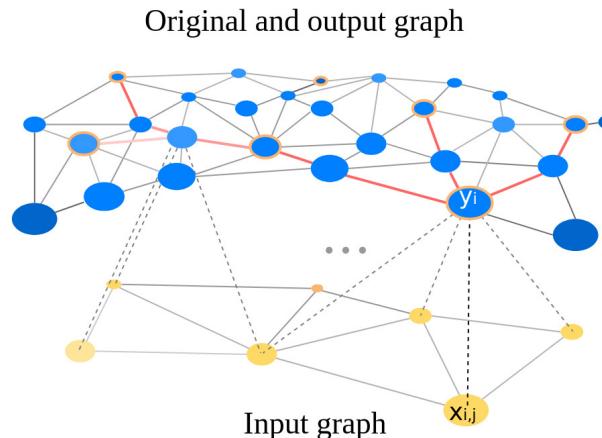
TransConv/Unpool, radius=2



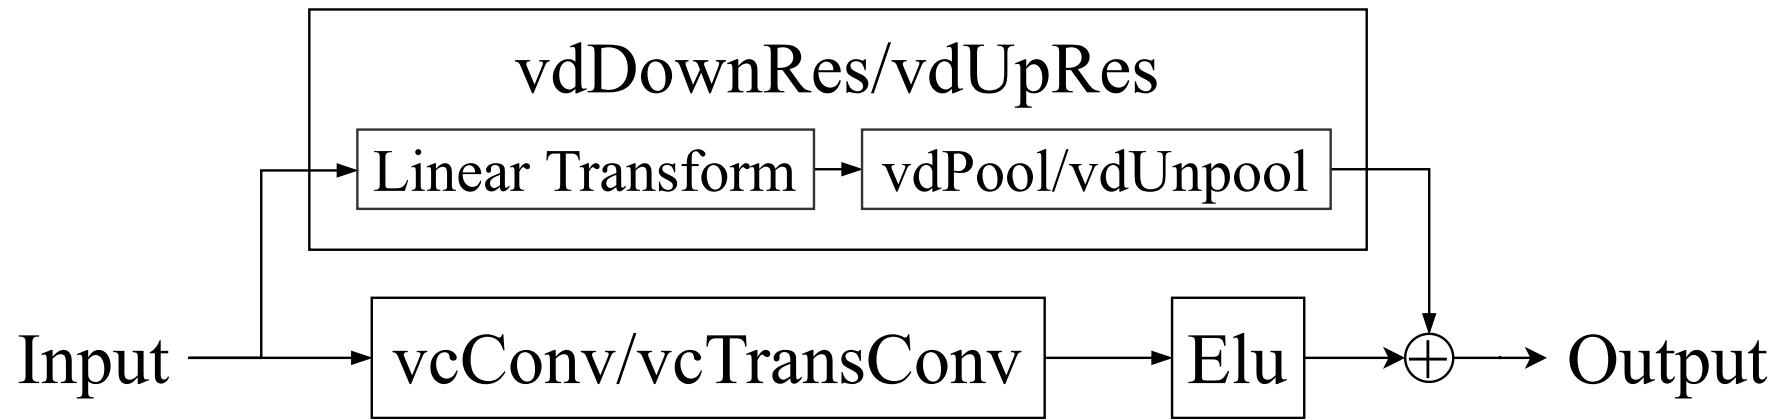
Create edges between input and output graphs.

Operations Analog to Regular CNN

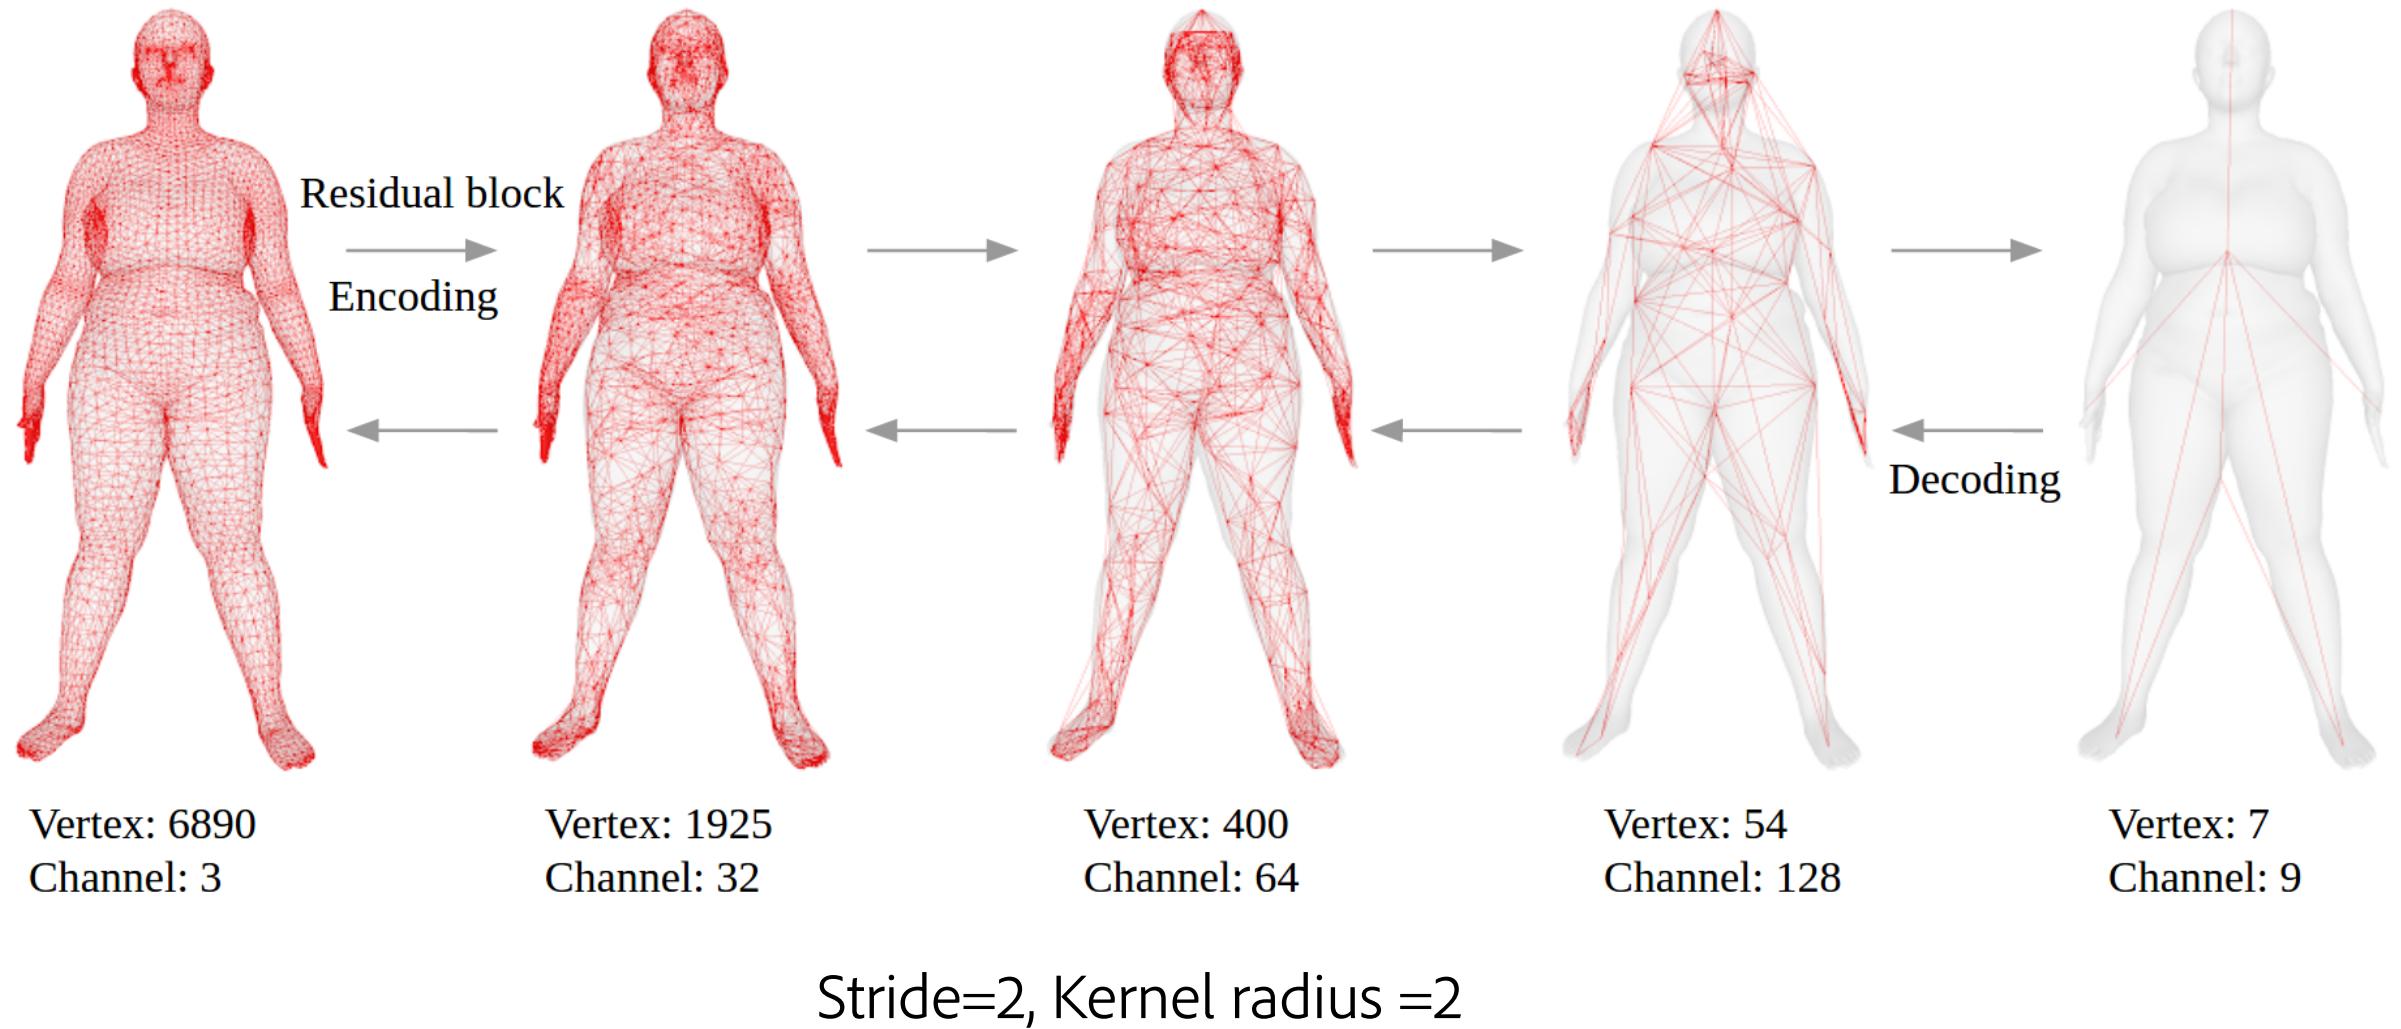
Down-sampling	Up-sampling	Attributes
vcConv	vcTransposeConv	(Stride, kernel radius, basis size, in_channel, out_channel, dilation)
vdPool	vdUnpool	(Stride)
vdDownResidual	vdUpResidual	(In_channel, out_channel)



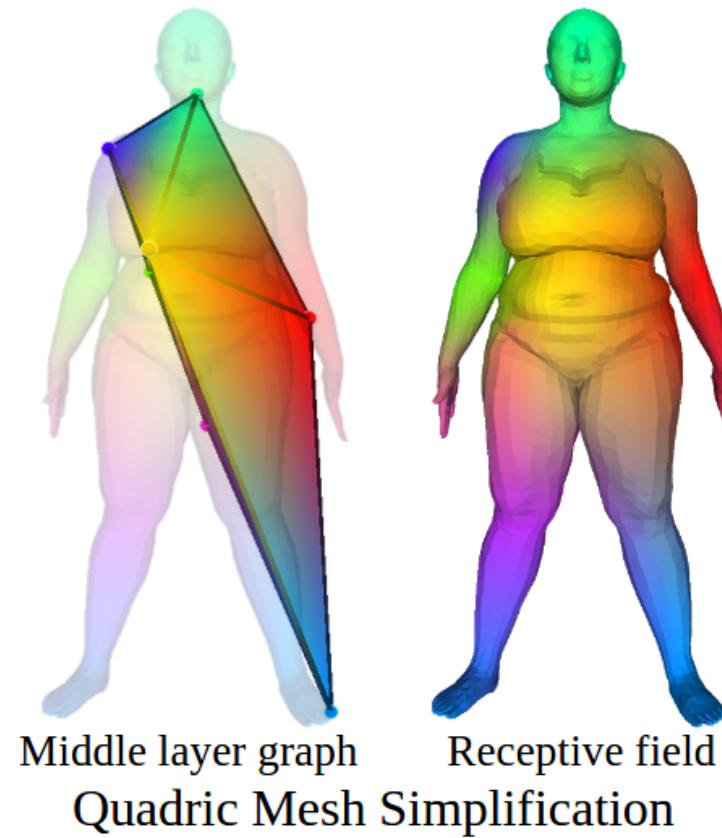
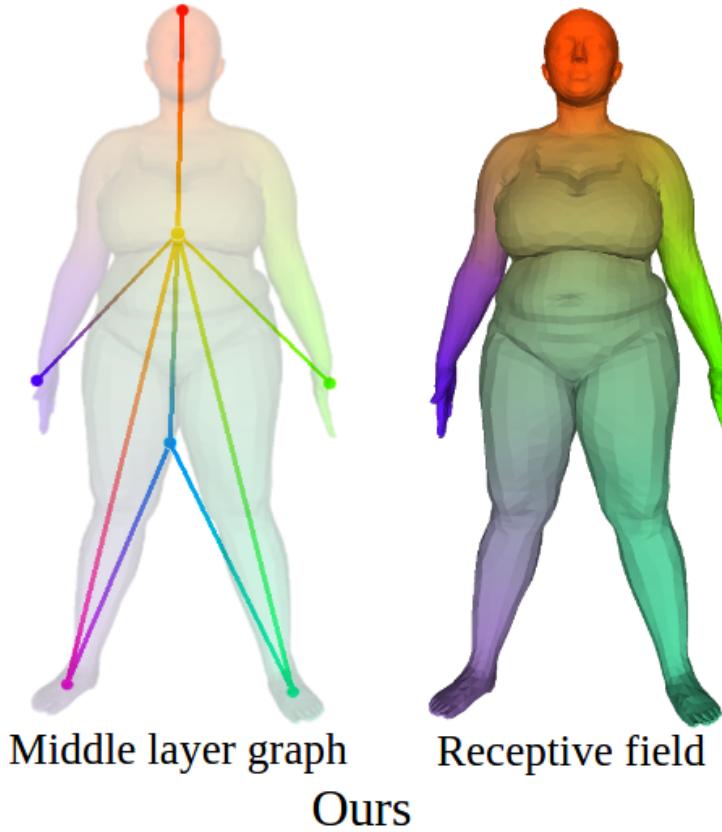
Residual Block



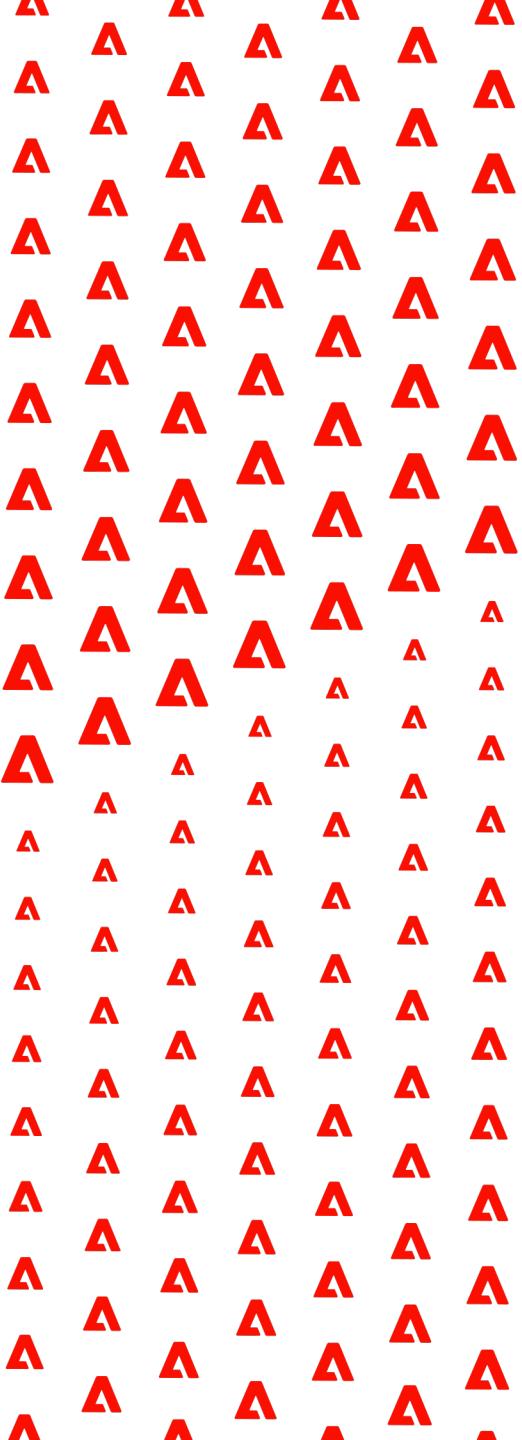
Auto-Encoder for DFAUST Dataset



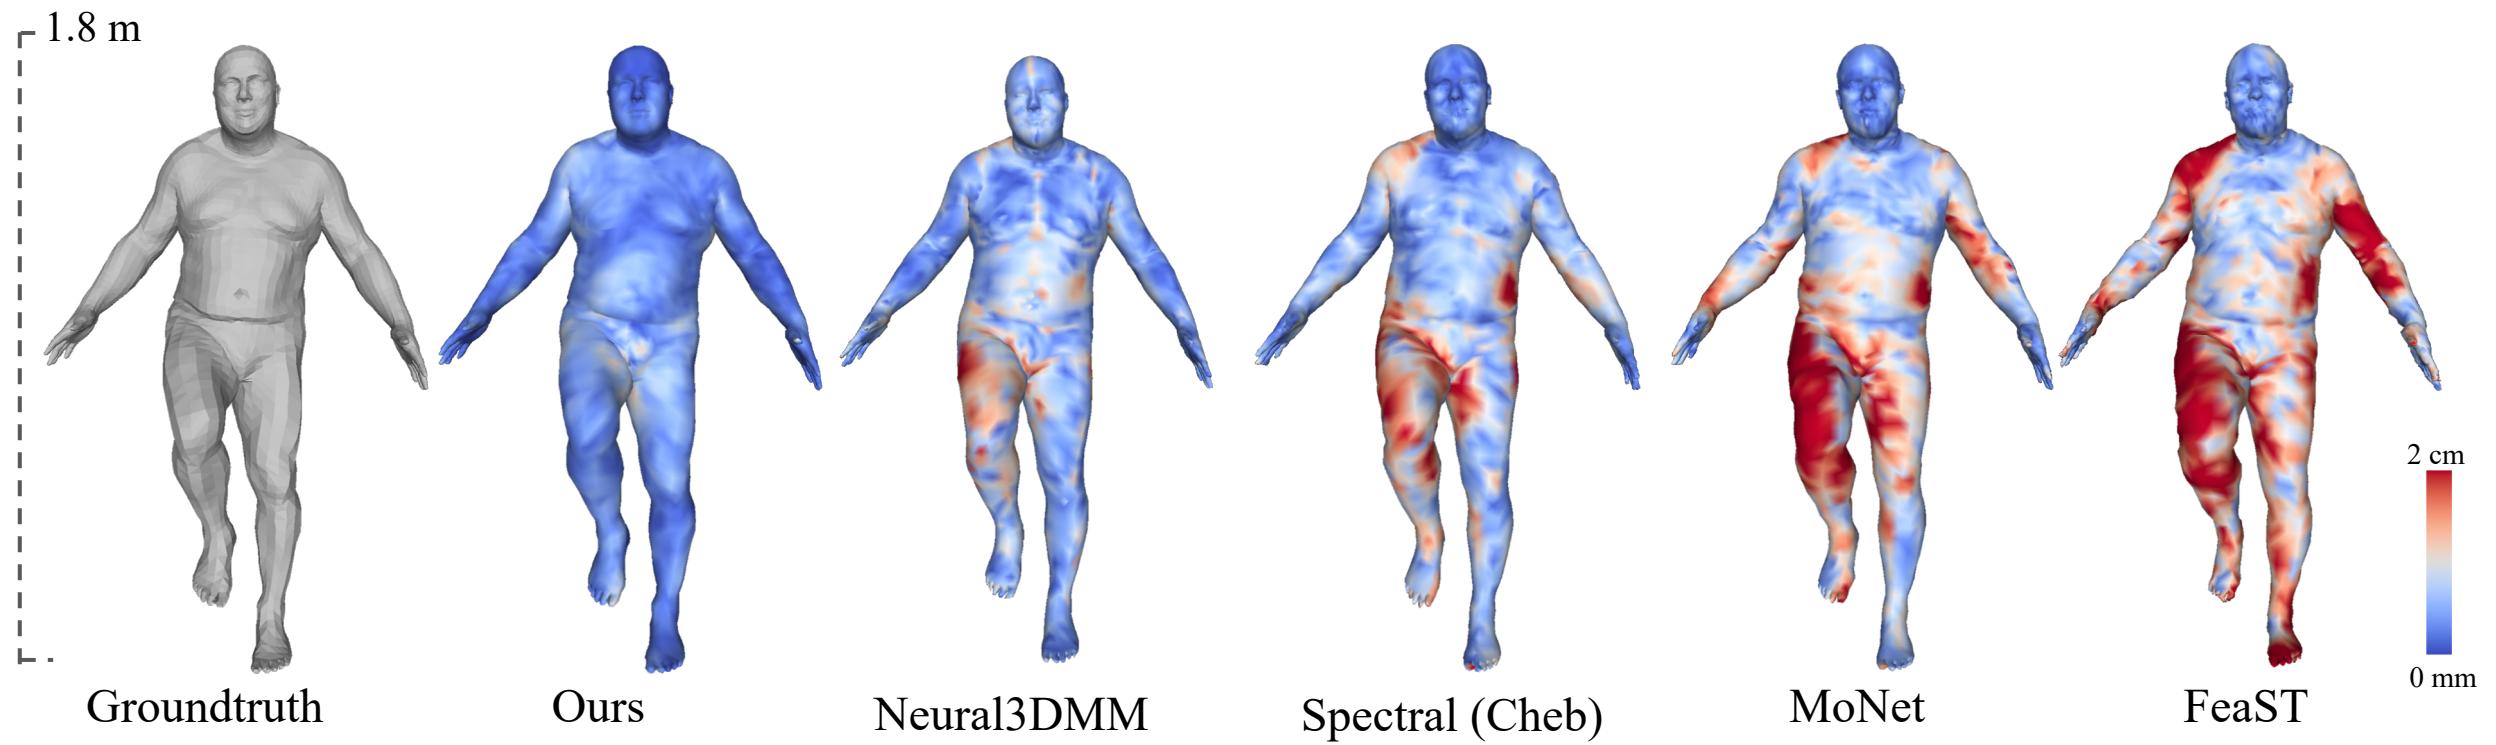
Localized Latent Features



Results

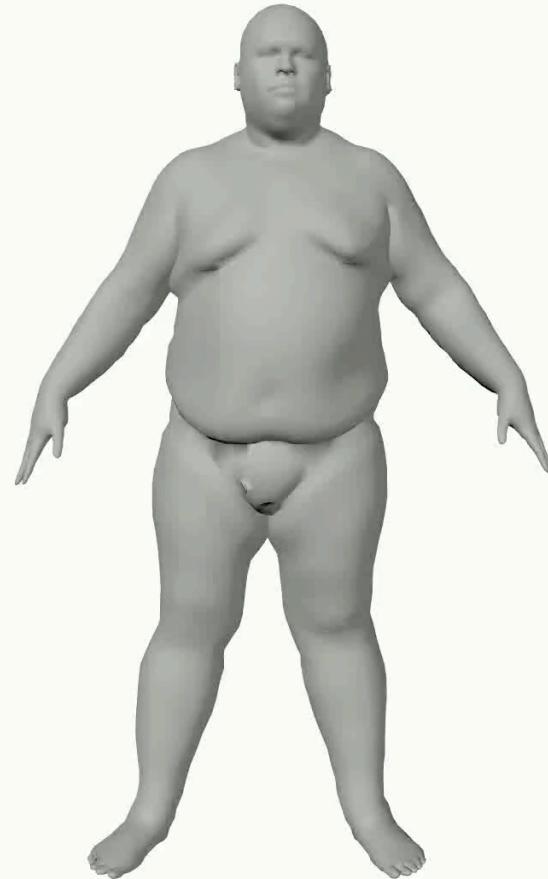


Ours: Lowest Reconstruction Error



Ours: Lowest Reconstruction Error

Compression Rate: 0.3%



Groundtruth



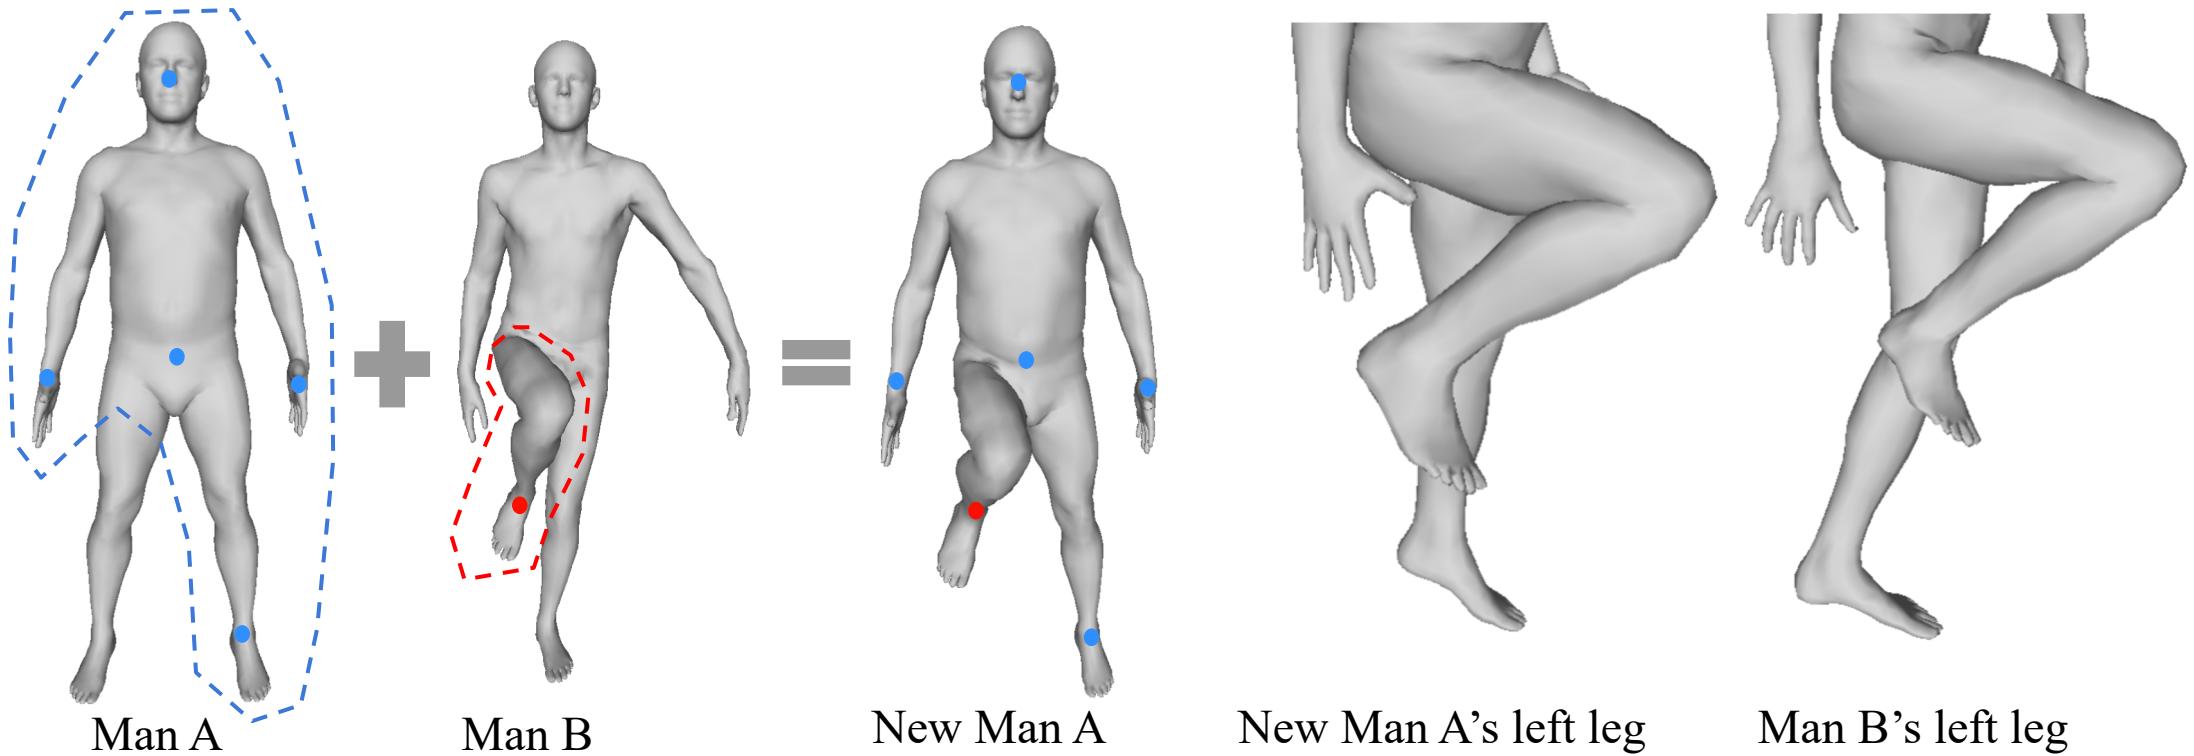
Ours



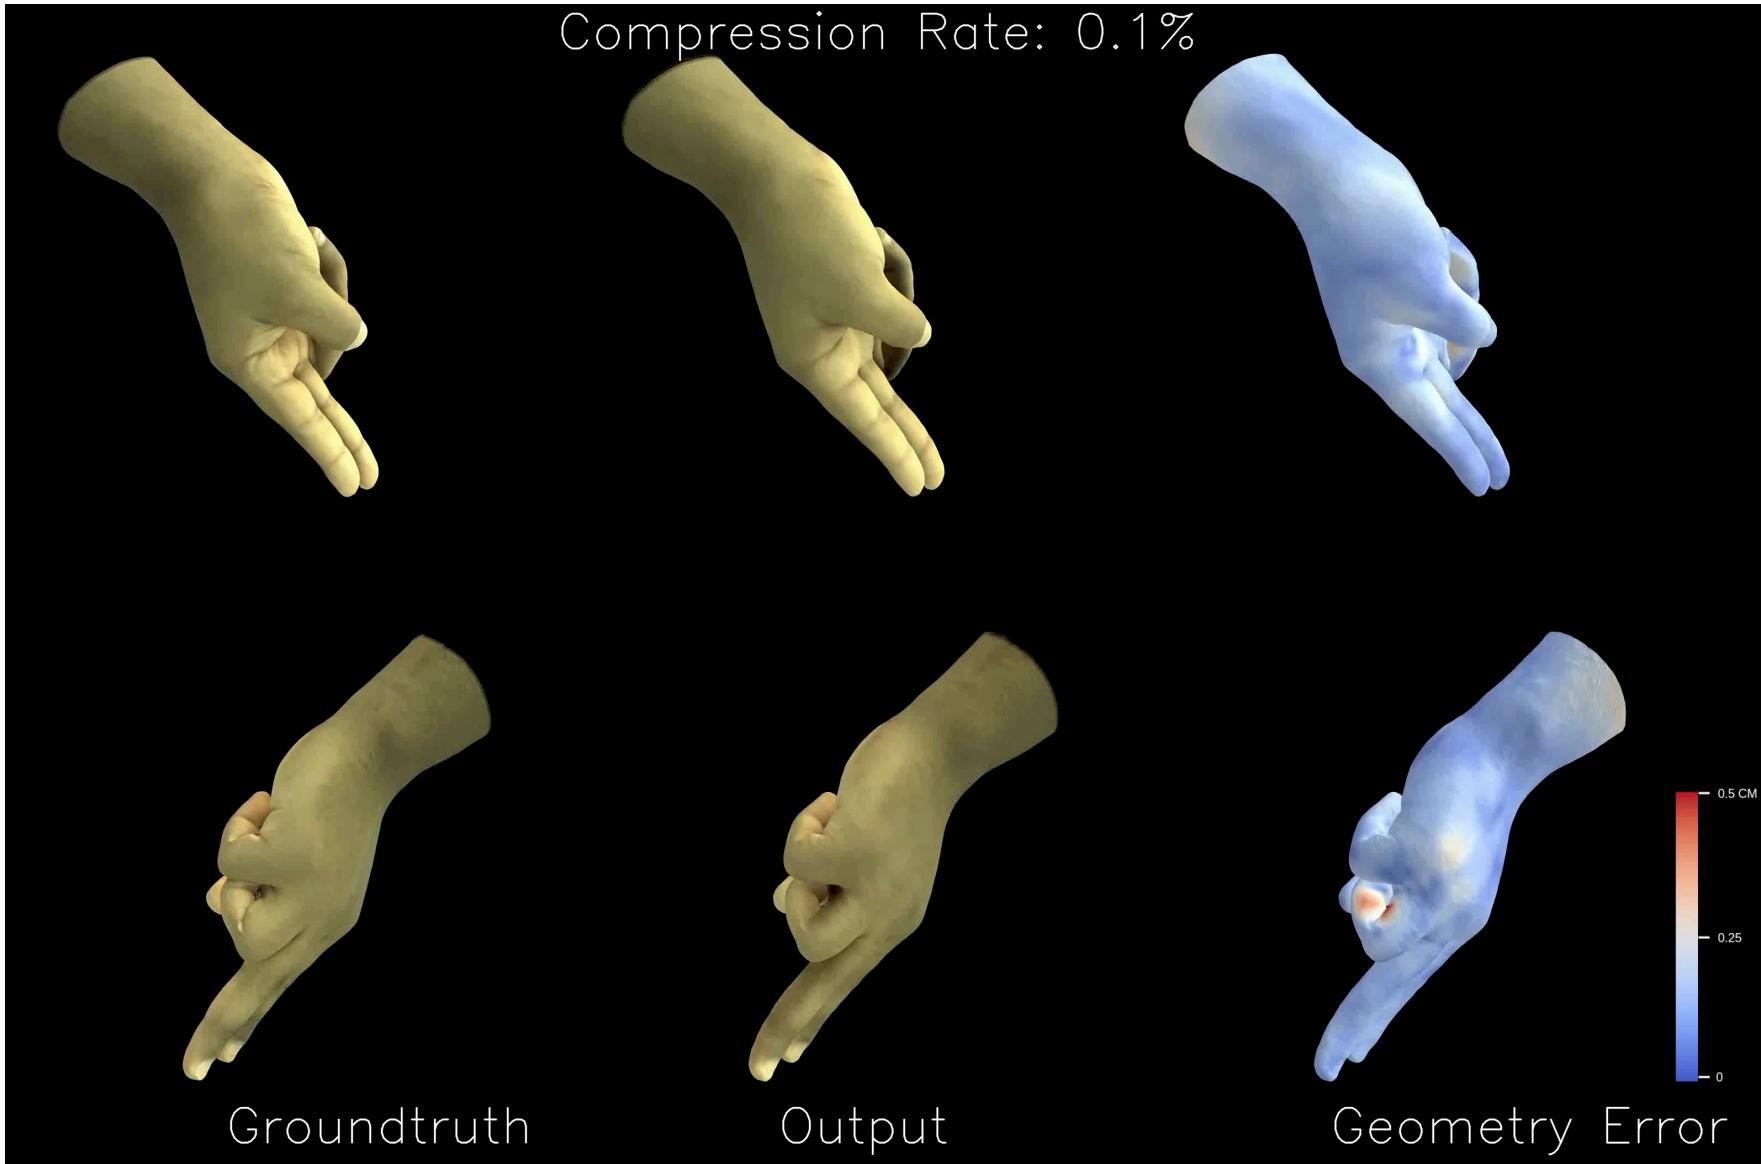
Nerual3DMM

Videos can be watched at https://zhouyisjtu.github.io/project_vcmeshcnn/vcmeshcnn.html

Localized Latent Space Interpolation

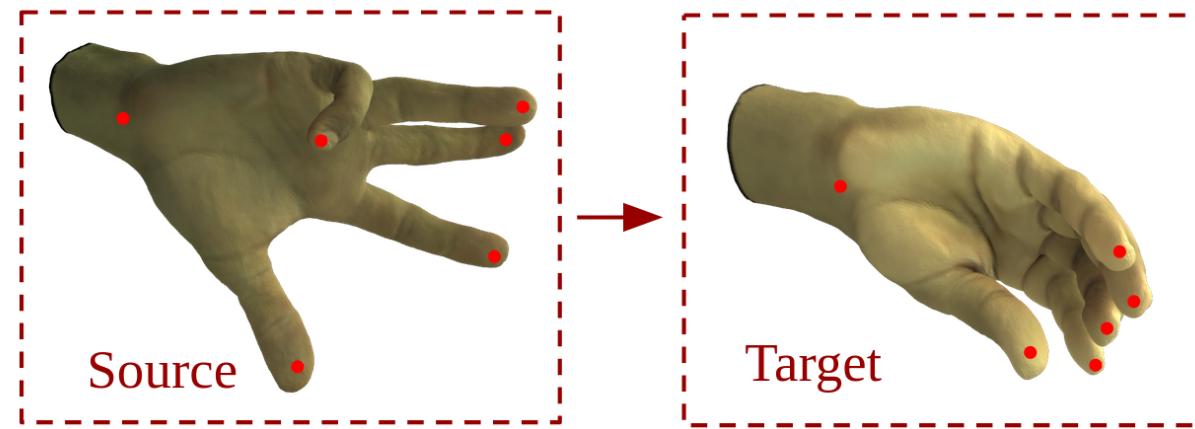
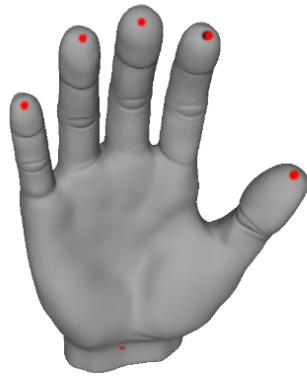


Reconstruct both Geometry and Color

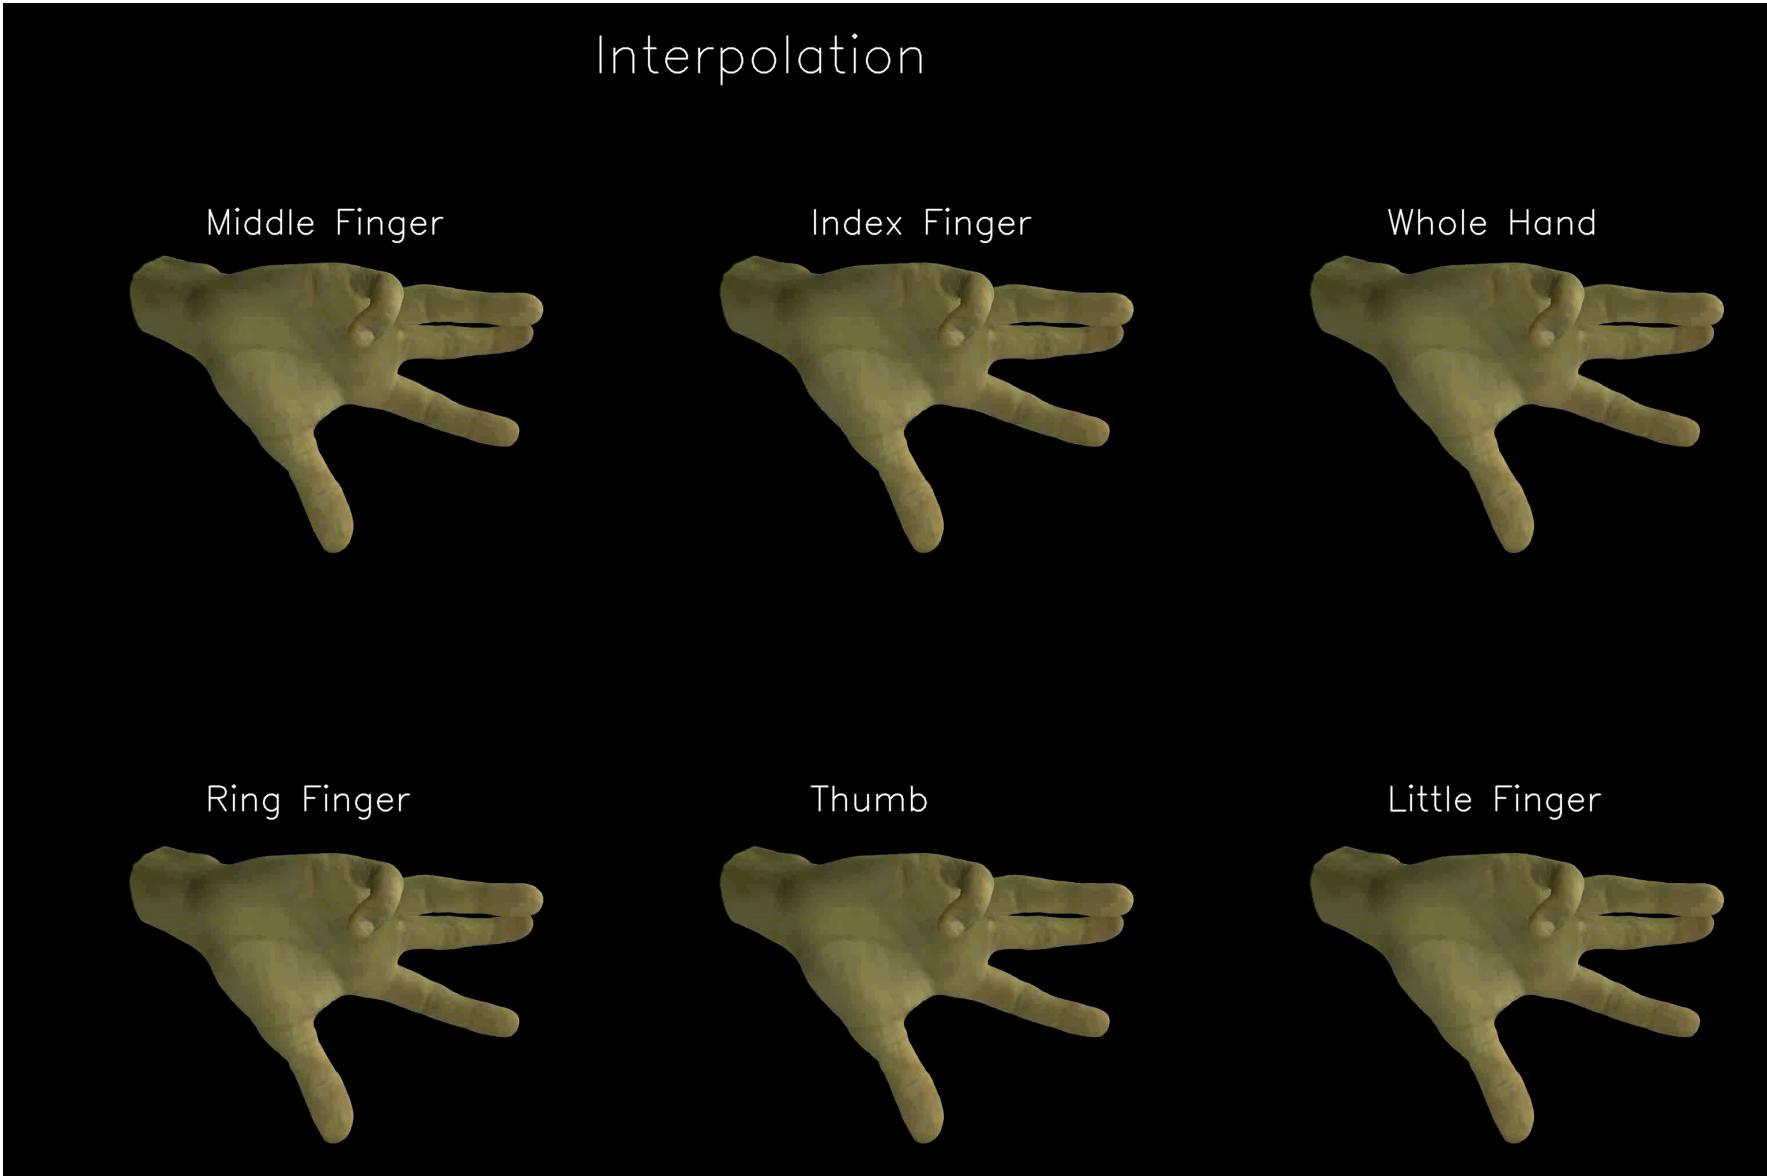


Localized Latent Space Interpolation

- Latent vertices

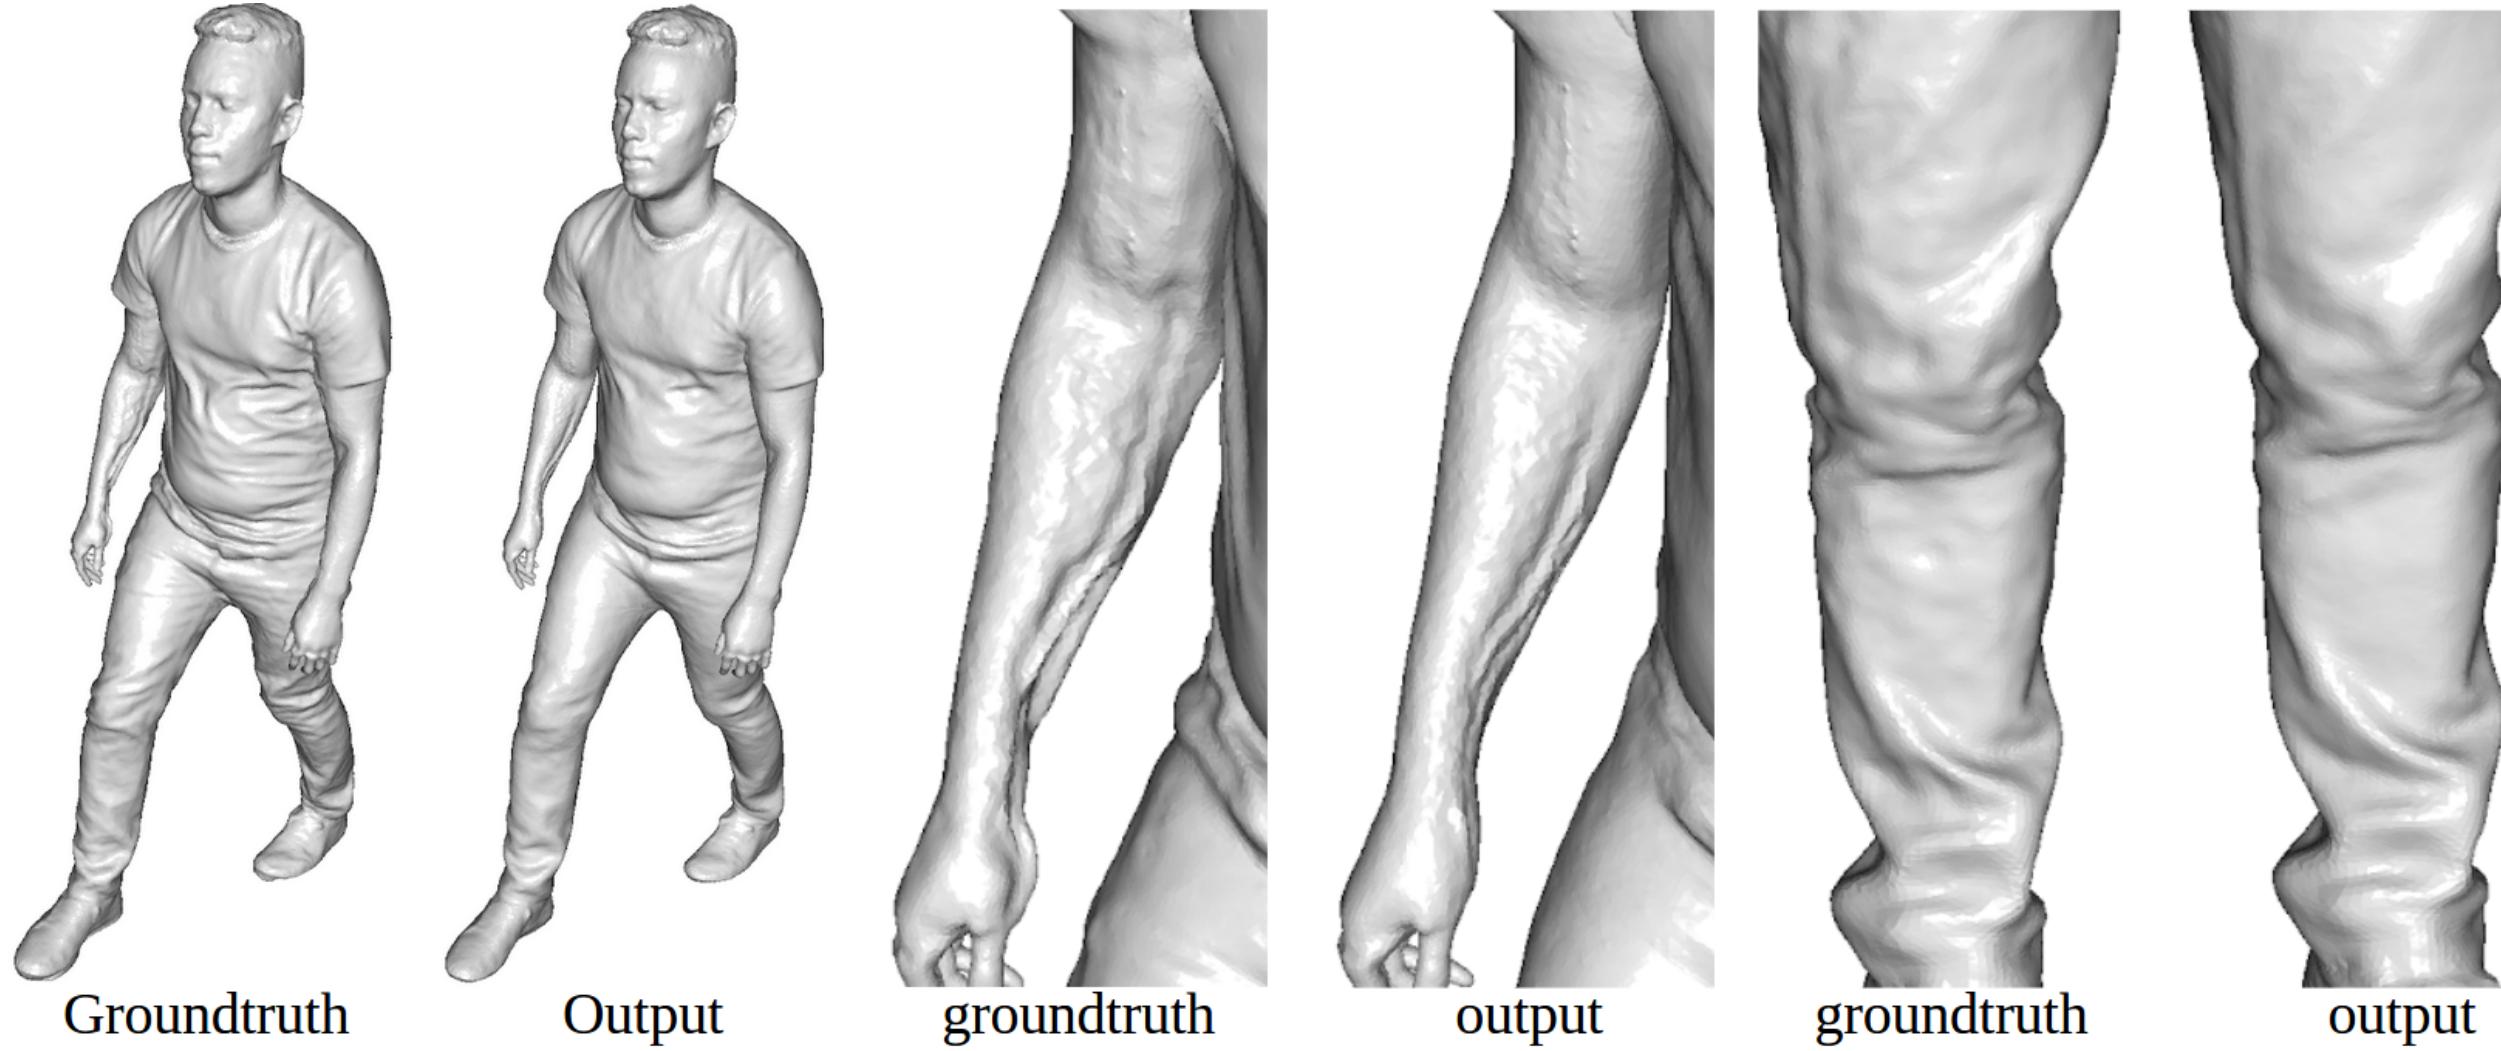


Localized Latent Space Interpolation



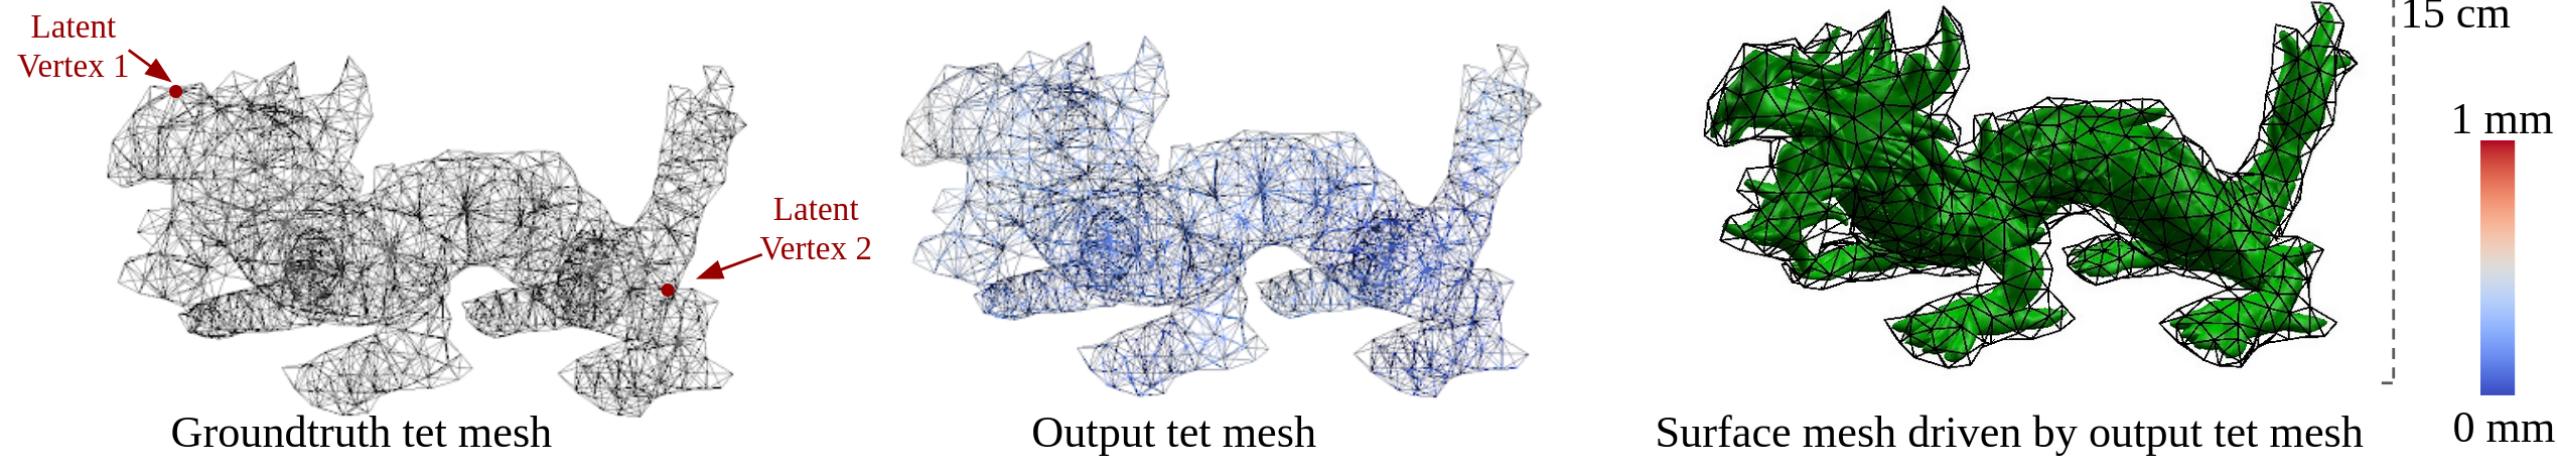
Efficient for High Resolution Meshes

153,000 Vertices, 24k training meshes, 2k test meshes, compression rate 0.75%.



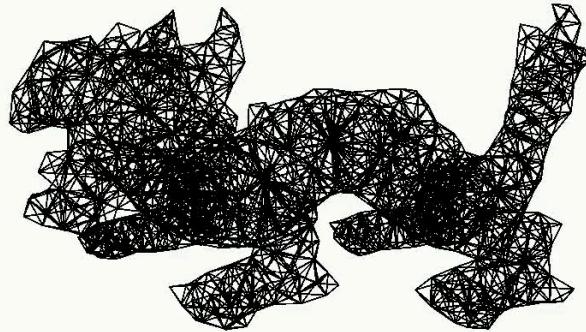
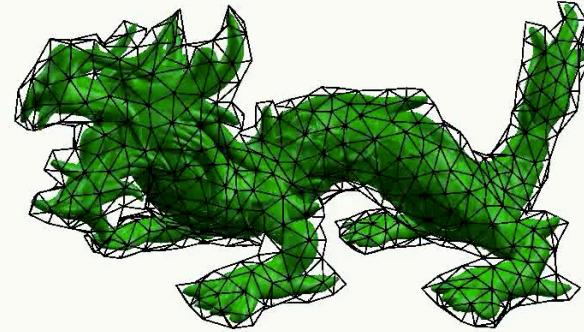
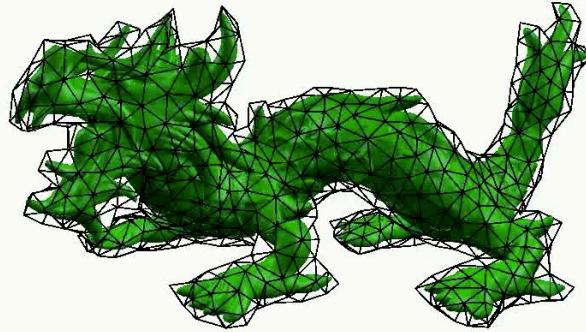
Volumetric Mesh (Tetrahedrons)

960 Vertices, 7k training meshes, 562 test meshes, compression rate 1.1%, test error 0.2 mm.

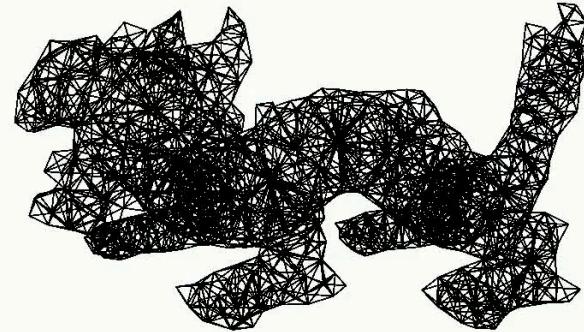


Volumetric Mesh (Tetrahedrons)

Compression Rate: 1.1%



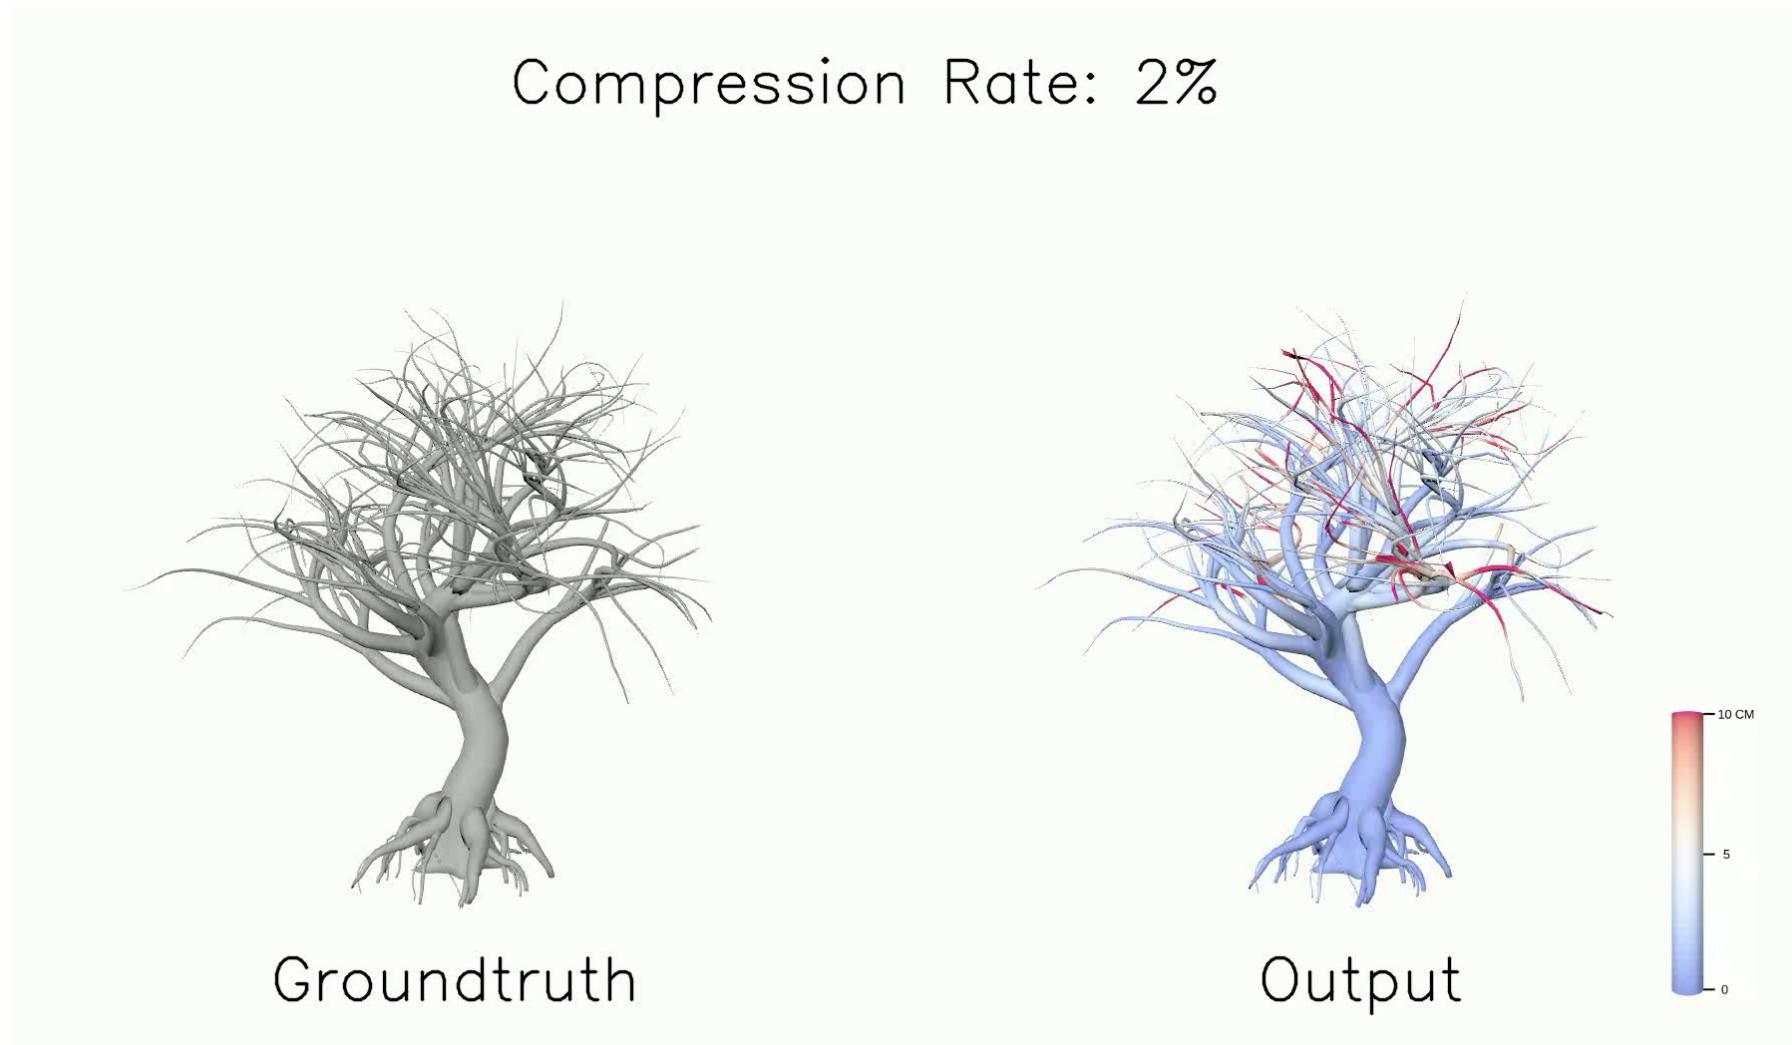
Groundtruth



Output

Non-Manifold Mesh

20,000 Vertices, 10k training meshes, 2k test meshes, compression rate 2%, test error 4.1 cm.



Future Work

